

8. Hash Tables

- Motivation
- Hash Tables
- Integer Hashing
- String Hashing
- Storage and Retrieval
- Hashing Collisions
- Linear Probing
- Double Hashing
- Separate Chaining
- Hash Buckets
- Hashing Analysis
- Hashing Discussion

Motivation

- So far we have seen a variety of ADTs for storing and retrieving data.
- Goal of Hash Tables is to get faster access!

ADT	Search Time
array	$O(N)$ – unsorted $O(\log N)$ – sorted
linked list	$O(N)$ – sorted or unsorted
stack	$O(1)$ – for top only
queue	$O(1)$ – for head/tail only
binary tree	$O(\log N)$ – for any value
heap	$O(\log N)$ – for largest value

Hash Tables

- A hash table is a data structure invented for very fast data storage and retrieval.
- Goal is to look for data in only ONE step using the data itself to tell you where to look.
- We use a hash function to map data values into table positions.

eg: $\text{hash}(\text{"john"}) = 42$ so we store "john" in position 42.


			...	john	...		
0	1	2		42			

- Hash functions are many to one, so we need an algorithm to resolve collisions (where 2 values map to the same table position).

Integer Hashing

- Assume that we want to store/retrieve 100 student records by their ID.
- We could allocate an array 1,000,000 long use ID as the index, but this would waste a lot of space.
- Instead, we can allocate an array 200 long and use $ID \% 200$ as our hash function.

eg: $\text{hash}(672403) = 3$



- We need to make sure the hash table size is larger than number of records we intend to store.

String Hashing

- Goal is to spread index values uniformly around the hash table to reduce collisions.
 - There are many ways to do this with a string.
 - add the ASCII codes for all characters in the string.
 - convert string into a number base 256.
 - select k characters from string and multiply ASCII codes by user defined position weights.
- eg. $\text{index} = (\text{str}[2]*17 + \text{str}[3]*21 + \text{str}[6]*33) \% \text{size}$

j	o	h	n		g	a	u	c	h
0	1	2	3	4	5	6	7	8	9

Storage and Retrieval


- To store a record in hash table:
 - calculate index.
 - check array[index] is empty.
 - store record in array[index].
 - mark location as “taken”.
- To retrieve a record from hash table:
 - calculate index.
 - check array[index] to see if not empty and key matches.
 - return record from array[index] or a “not found” message.
- This approach is very fast but there is one potential problem ...

Hashing Collisions

- When we attempt to store a value and the location is already taken this is called a collision.
- Instead of giving up, we must store the data somewhere else in the hash table.
- Many collision resolutions options are possible:
 - Linear probing.
 - Double hashing.
 - Separate chaining.
 - Hash buckets.
- We also need to modify our retrieval algorithm to look in alternative locations if necessary.

Linear Probing


- If the location at index is “taken” simply probe the next locations until an open spot is located.



- Here we use -1 to mark an open spot in the table.
- New data replaces the -1 and now that location is “taken”.
- We must adapt our lookup to probe until -1 reached to check the locations adjacent to hash table index.
- If data is deleted, we must mark with “deleted” flag and search/insert adjusted accordingly.

Double Hashing

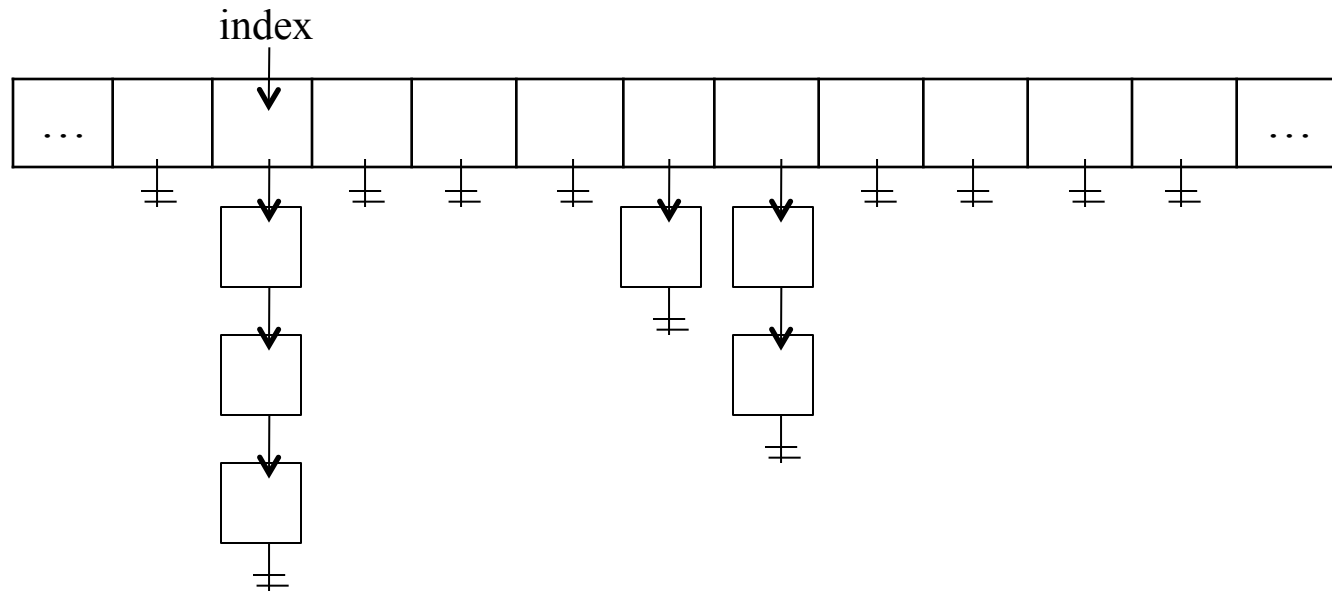
- Similar idea to linear probing except the step size between probes is a function of the hash index.
eg: $\text{step} = (\text{index} + 79) \% 23 + 1$
- This will spread out the data records in case of collisions.
- Probes should wrap around at the end of the table (using the modulo % operator).
- Double hashing works best if the step size is relatively prime to table size. Hence use a prime table size.



- Here, $\text{index} = (\text{index} + \text{step}) \% \text{size}$

Separate Chaining


- Instead of probing the hash table we can use a linked list or another dynamic ADT to store collisions.



- Insertion/search/deletion now use our existing linked list code
- Since each list is very small, the operations take very few steps.
- The table size for separate chaining can be smaller because the overflow space is in the linked list.

Hash Buckets

- We can also use fixed size buckets to resolve collisions instead of dynamic data structures.



- The hashing function returns the index of the first location in each bucket.
- We use linear probing to locate empty location in bucket to store data.
- If the bucket becomes full, then we probe next bucket to find empty spot.
- Same space requirements as linear probing and double hashing.

Hashing Analysis

- Speed of insert/search/delete depends on the number of collisions
- Let α between 0 and 1 be fraction of table that is occupied, hence $1 - \alpha$ is probability location is empty.

Probe	Occupied	Free
1	α	$(1-\alpha)$
2	α^2	$(1-\alpha) \alpha$
3	α^3	$(1-\alpha) \alpha^2$
4	α^4	$(1-\alpha) \alpha^3$
...
n	α^n	$(1-\alpha) \alpha^{(n-1)}$

- To calculate the average number of probes to locate a free location we sum the product of probe and free columns

Hashing Analysis

$$S = 1(1 - \alpha) + 2(1 - \alpha)\alpha + 3(1 - \alpha)\alpha^2 + \dots n(1 - \alpha)\alpha^{n-1}$$
$$\alpha S = 1(1 - \alpha)\alpha + 2(1 - \alpha)\alpha^2 + \dots (n - 1)(1 - \alpha)\alpha^{n-1}$$

Subtracting and simplifying:

$$(1 - \alpha)S = (1 - \alpha) + (1 - \alpha)\alpha + (1 - \alpha)\alpha^2 + \dots (1 - \alpha)\alpha^{n-1}$$

$$S = 1 + \alpha + \alpha^2 + \dots \alpha^{n-1} = \frac{1}{1 - \alpha}$$

examples:

$$\alpha = 1/4, S = 4/3$$

$$\alpha = 1/2, S = 2$$

$$\alpha = 3/4, S = 4$$

$$\alpha = 1/10, S = 10/9$$

- Rule of thumb: make hash tables 2-4 times larger than amount of data you expect to store to keep number of probes small.

Hashing Discussion

- If we wish to store N values in a hash table and we allocate a table of size $K*N$ then the highest value possible for $\alpha=1/K$.
- The average number of probes to find a free location will be:

$$S = \frac{1}{1 - \alpha} = \frac{1}{1 - 1/K} = \frac{K}{K - 1}$$

- S does not depend on N , so hashing runs in constant time $O(1)$
- This is clearly much better than our other ADTs where access time is $O(\log N)$ at best.
- Hash tables are widely used in applications when the data can fit in memory.
- Binary search trees are a better choice when data sets are too large for memory and must be stored on disk.